BiomedCLIP finetuning on GradCAM heatmaps

training (finetune.py, FLYP-style)

data/ : 12,890 GradCAM heatmaps per factor (z0 = intensity, z1 = area), split 90/10 by cell | batch 64, 5 epochs

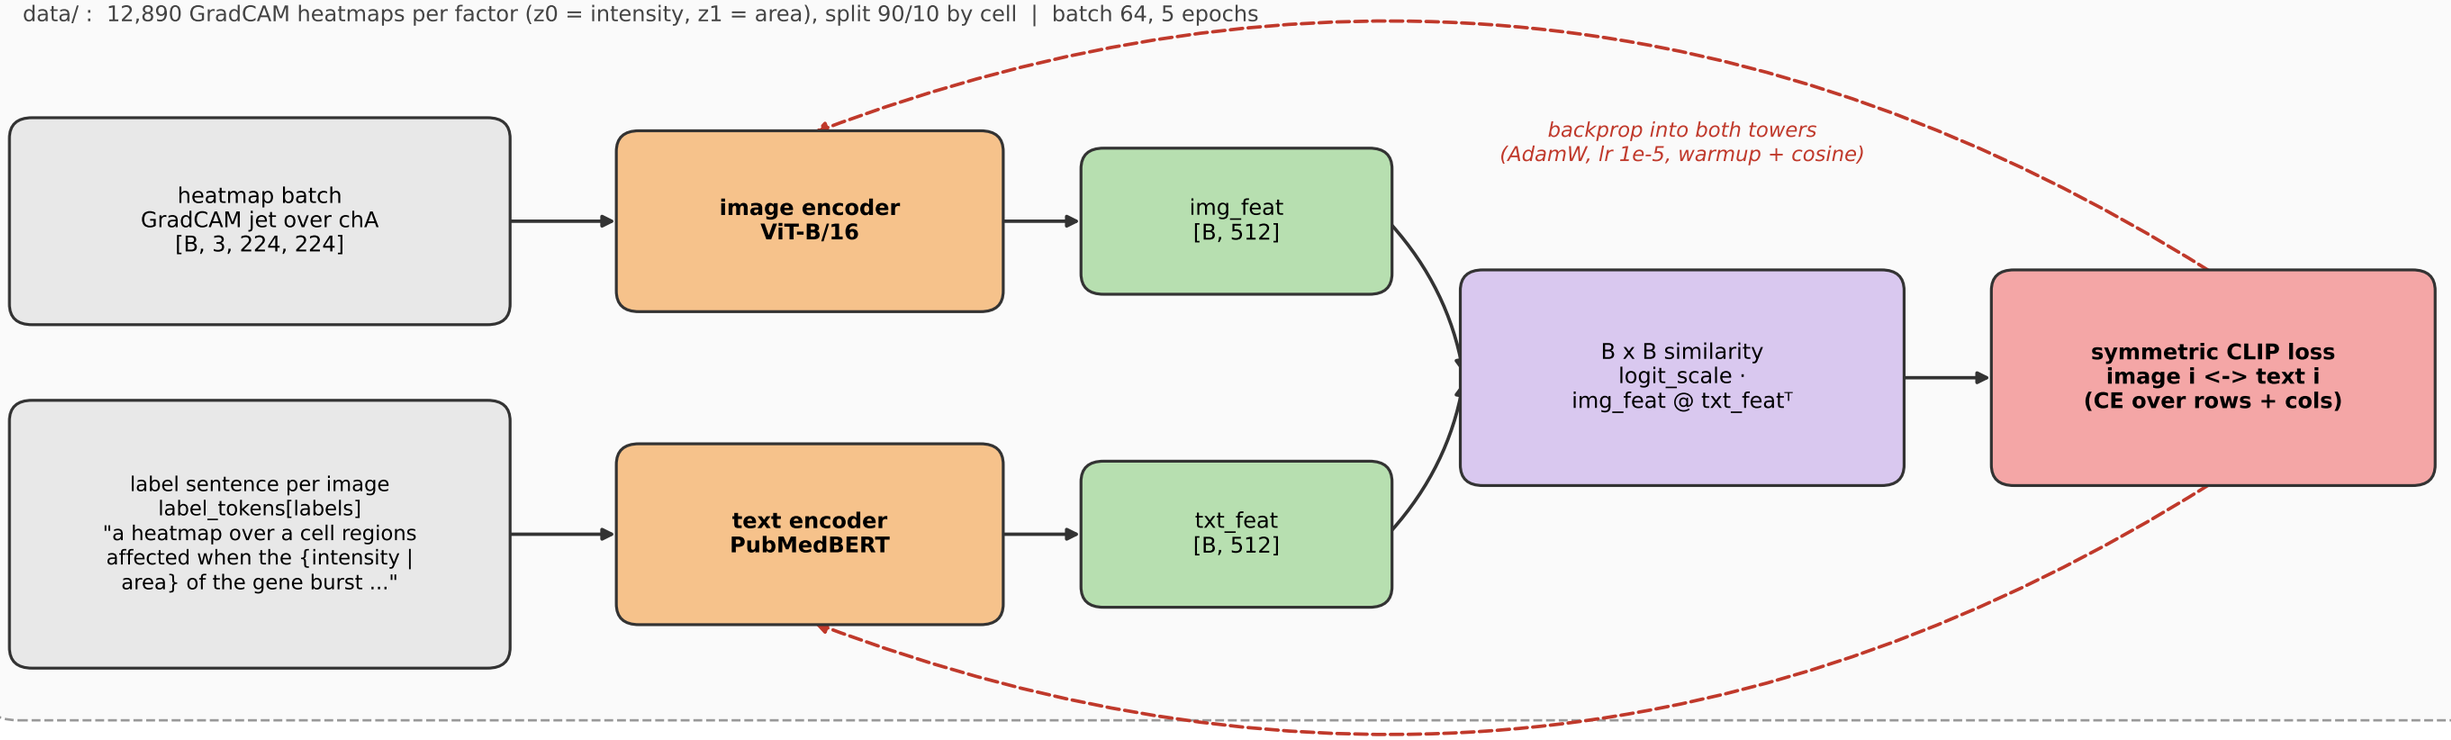

evaluation (finetune.evaluate per epoch, eval.py per heatmap)

val split of data/ (seen heatmaps), or data_unknown/z2..z15 to probe the unseen labels | best.pt = best val seen acc

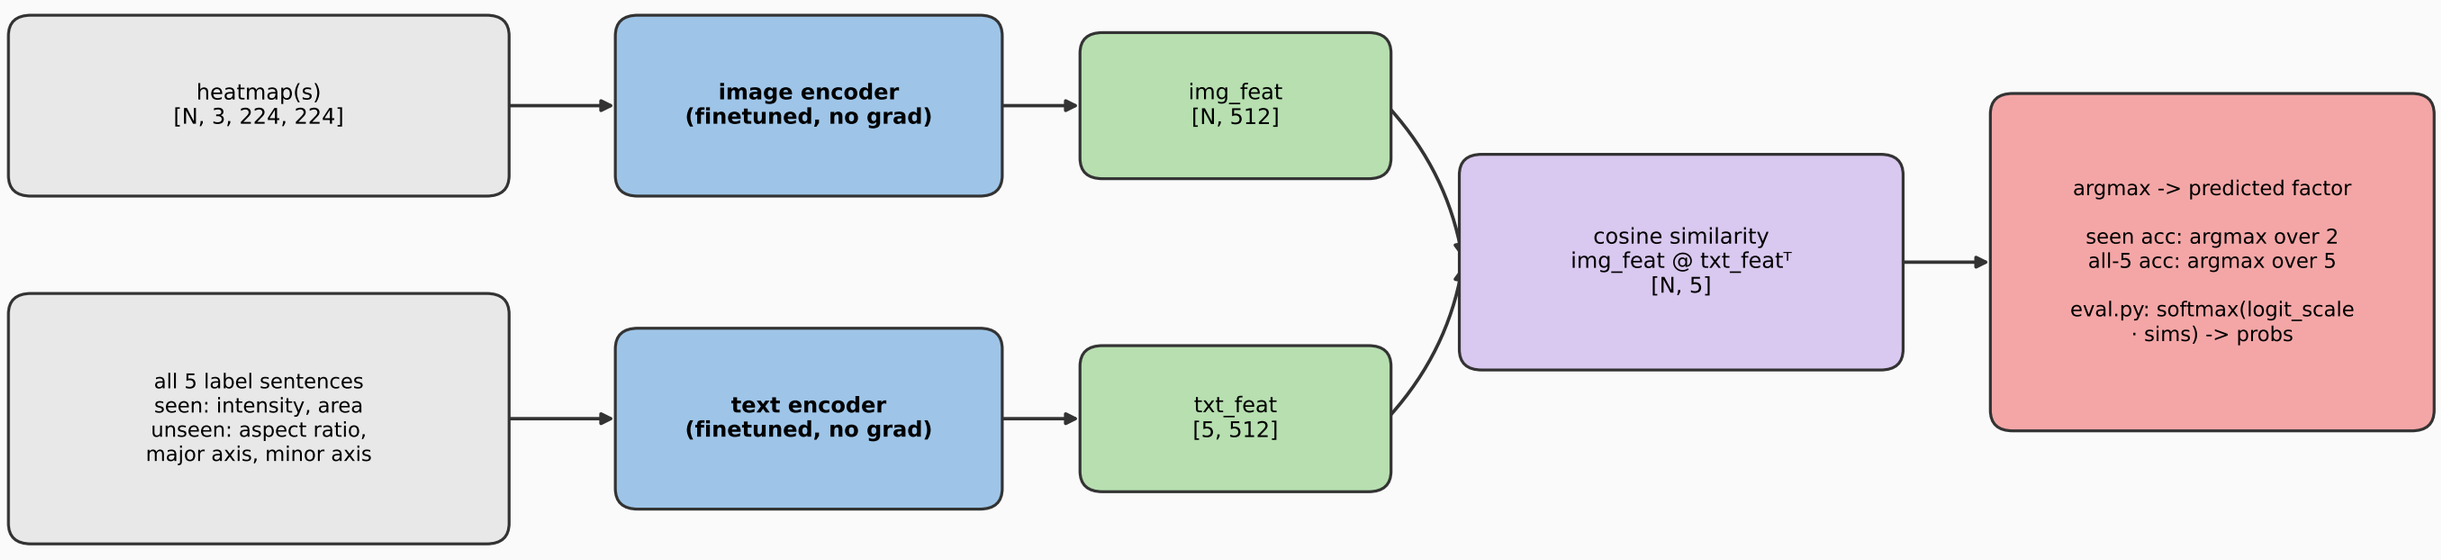

input | trained | finetuned, fixed | features | similarity | loss / output